

Homework 3

Yuanrui Zhang (yz545)

Problem 1. Programming Models

(1) “for i” loop only

- Read-only: N, data_array
- Read/write non-conflicting: data_gridX, data_gridY
- Read/write conflicting: i, j, sum, product, measurement

(2) “for j” loop only

- Read-only: N, data_array, sum, i
- Read/write non-conflicting: None
- Read/write conflicting: j, data_gridX, data_gridY, product, measurement.

Problem 2. Code Analysis for Parallel Task Identification


Loop-carried dependency:

$$(1) S1[i, j] \rightarrow T S1[i, j - 2]$$

$$(2) S2[i, j] \rightarrow T S2[i + 1, j]$$

$$(3) S2[i, j] \rightarrow A S2[i + 1, j]$$

LDG:



- (a) No, it isn't. Because of the loop carried dependency (2) and (3) list above, there can't be DOALL parallelism.
- (b) No, the same reason as (a) – there is a loop carried dependency (1) between for j loop iterations.
- (c) The update of each node along a diagonal is an independent parallel task because there is chain of dependency in this pattern, e.g., $[1, 3] \rightarrow [1, 1] \rightarrow [2, 1] \rightarrow [3, 1]$, but along anti-diagonal is not.

(d) Yes. We can also parallelize for j loop with odd indices and with even indices. Because according to the LDG, the loop-carried dependency only exists among odd columns or among even columns. That's said, no such dependency exists across an even and an odd column. Thus, we can effectively divide it into 2 parallel tasks.

Problem 3. Code Profiling & Performance Counters

(a) Performance profiling

Function index	Function name	The number of calls	percentage of execution time
1	miniFE::matvec_std<miniFE::CSRMatrix<double, int, int>, miniFE::Vector<double, int, int> >::operator()(miniFE::CSRMatrix<double, int, int>&, miniFE::Vector<double, int, int>&, miniFE::Vector<double, int, int>&)	201	30.09%
2	frame_dummy	1597918831	11.08%
3	std::_Rb_tree<int, int, std::_Identity<int>, std::less<int>, std::allocator<int> >::_S_key(std::_Rb_tree_node<int> const*)	57598102	6.52%
4	std::_Rb_tree<int, int, std::_Identity<int>, std::less<int>, std::allocator<int> >::_S_value(std::_Rb_tree_node<int> const*)	435792686	4.48%
5	int* std::lower_bound<int*, unsigned long>(int*, int*, unsigned long const&)	32768000	3.69%
6	__gnu_cxx::__aligned_membuf<int>::_M_ptr()	435883250	3.22%
7	std::_Rb_tree_node<int>::_M_valptr()	435883250	2.99%
8	miniFE::decide_how_to_shrink(Box const&, Box const&)	719	2.99%

(b) Amdahl's Law

$$\text{Speedup} = 1 / (1 - 0.3009 + 0.3009 / 5) = 1.32$$

(c) Performance Counters

(1) When running `perf stat ./miniFE.x -nx 40 -ny 80 -nz 160` by default:

```
Performance counter stats for './miniFE.x -nx 40 -ny 80 -nz 160':

120881.527723      task-clock (msec)    #    1.000 CPUs utilized
          36      context-switches        #    0.000 K/sec
           0      cpu-migrations          #    0.000 K/sec
       51,327      page-faults            #    0.425 K/sec
284,563,542,170    cycles                  #    2.354 GHz
124,804,275,916    stalled-cycles-frontend #   43.86% frontend cycles idle
<not supported>   stalled-cycles-backend
433,569,432,522    instructions          #    1.52 insns per cycle
                                     #    0.29 stalled cycles per insn
 88,728,012,007    branches                # 734.008 M/sec
 2,044,921,361     branch-misses          #    2.30% of all branches

120.891166326 seconds time elapsed
```

- (2) Measure Instructions; CPU cycles (and also show IPC, instructions per cycle); Branch instructions; Branches misses (mispredictions)

```
Performance counter stats for './miniFE.x -nx 40 -ny 80 -nz 160':

281,371,002,033      cpu-cycles
433,564,078,692      instructions          #    1.54  insns per cycle
88,726,944,642       branches
1,872,806,757        branch-misses          #    2.11% of all branches

119.743417490 seconds time elapsed
```

- (3) Measure Cache references; L1 data cache load misses; L1 instruction cache load misses; LLC (last level cache) loads; LLC (last level cache) load misses; Data TLB load misses

```
Performance counter stats for './miniFE.x -nx 40 -ny 80 -nz 160':

338,645,856          cache-references              (66.67%)
767,371,473          L1-dcache-load-misses          (66.67%)
11,115,297           L1-icache-load-misses          (66.67%)
655,282,132          LLC-loads              (66.67%)
350,068,628          LLC-load-misses          #  53.42% of all LL-cache hits (66.67%)
11,679,454           dTLB-load-misses              (66.66%)

120.129659673 seconds time elapsed
```

Problem 4. Performance Counters

1. Re-run my program across different loop nest orderings on the machine where I am using 'perf'

Loop nest orderings	Time
I-J-K	14.641327 s
I-K-J	0.716340 s
J-K-I	27.432645 s

2. Use 'perf' to see performance counters

(1) I-J-K

```
*****I-J-K*****
Time = 14.508862 s

Performance counter stats for './matrix 1':

    182,402,594      cache-references              (66.65%)
    2,161,328,049    L1-dcache-load-misses          (66.68%)
         625,832     L1-icache-load-misses          (66.69%)
    181,109,333      LLC-loads                  (66.69%)
     4,702,837       LLC-load-misses             #    2.60% of all LL-cache hits (66.68%)
    1,080,018,666     dTLB-load-misses            (66.65%)

    14.577152333 seconds time elapsed
```

(2) I-K-J

```
*****I-K-J*****
Time = 0.749057 s

Performance counter stats for './matrix 2':

    11,425,959      cache-references              (66.62%)
    139,369,377     L1-dcache-load-misses          (66.62%)
         91,449     L1-icache-load-misses          (66.62%)
    11,042,391      LLC-loads                  (66.98%)
         245,795     LLC-load-misses             #    2.23% of all LL-cache hits (67.02%)
     3,313,330      dTLB-load-misses            (66.66%)

    0.815300188 seconds time elapsed
```

(3) J-K-I

```
*****J-K-I*****
Time = 27.226312 s

Performance counter stats for './matrix 3':

    534,973,081      cache-references              (66.65%)
    3,228,790,274     L1-dcache-load-misses          (66.67%)
         957,443     L1-icache-load-misses          (66.68%)
    534,981,252      LLC-loads                  (66.68%)
         788,012     LLC-load-misses             #    0.15% of all LL-cache hits (66.67%)
    1,892,331,671     dTLB-load-misses            (66.65%)

    27.300640638 seconds time elapsed
```

From the three screenshots above, the first column of performance counter stats represents the raw counter numbers. We are easy to tell from those numbers that I-K-J has the least

number of all kinds of cache misses, while J-K-I has the most and I-J-K stays in between. In summary, performance counter results explain why these three patterns have different performance.