

# Homework4 Report

Yuanrui Zhang (yz545) & Yuchuan Li (yl645)

## Problem 1: Histogram

### Correctness:

```
yl645@vcm-17141:~/openmp/histo$ ./histo_locks uiuc.pgm > lock
yl645@vcm-17141:~/openmp/histo$ ./histo_atomic uiuc.pgm > atom
yl645@vcm-17141:~/openmp/histo$ ./histo_creative uiuc.pgm > creative
yl645@vcm-17141:~/openmp/histo$ diff lock validation.out
262,263d261
<
< Runtime =      1.99 seconds
yl645@vcm-17141:~/openmp/histo$ diff atom validation.out
262,263d261
<
< Runtime =      0.59 seconds
yl645@vcm-17141:~/openmp/histo$ diff creative validation.out
262,263d261
<
< Runtime =      0.33 seconds
```

The only difference between my results and the validation.out file is the runtime part, not the result of histogram part, so these three versions are all correct.

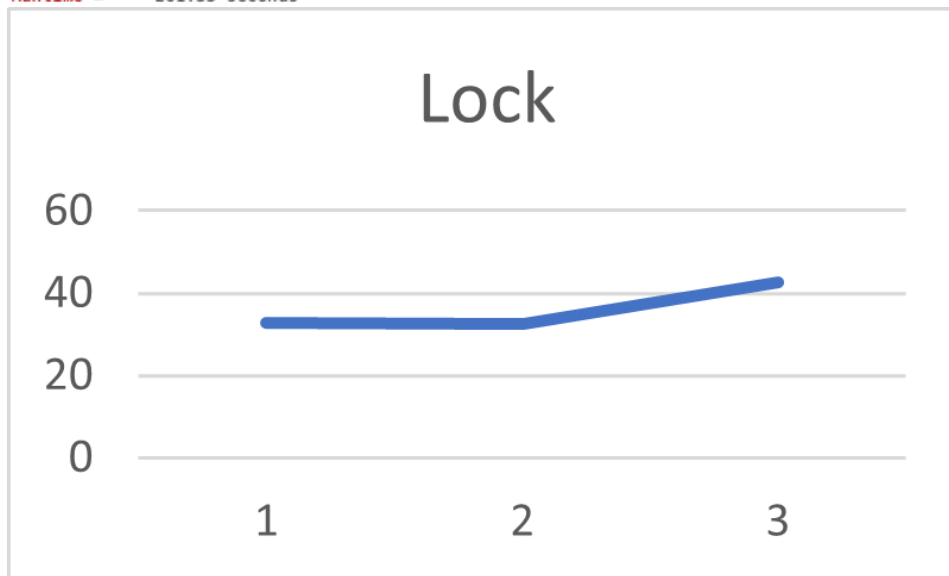
### Performance:

#### Original Version:

```
yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=2 ./histo uiuc-large.pgm | grep Runtime
Runtime =      6.20 seconds
yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=4 ./histo uiuc-large.pgm | grep Runtime
Runtime =      6.18 seconds
yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=8 ./histo uiuc-large.pgm | grep Runtime
Runtime =      6.15 seconds
```

#### Lock Version:

```
yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=2 ./histo_locks uiuc-large.pgm | grep Runtime
Runtime =     203.95 seconds
yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=4 ./histo_locks uiuc-large.pgm | grep Runtime
Runtime =     200.47 seconds
yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=8 ./histo_locks uiuc-large.pgm | grep Runtime
Runtime =     261.35 seconds
```



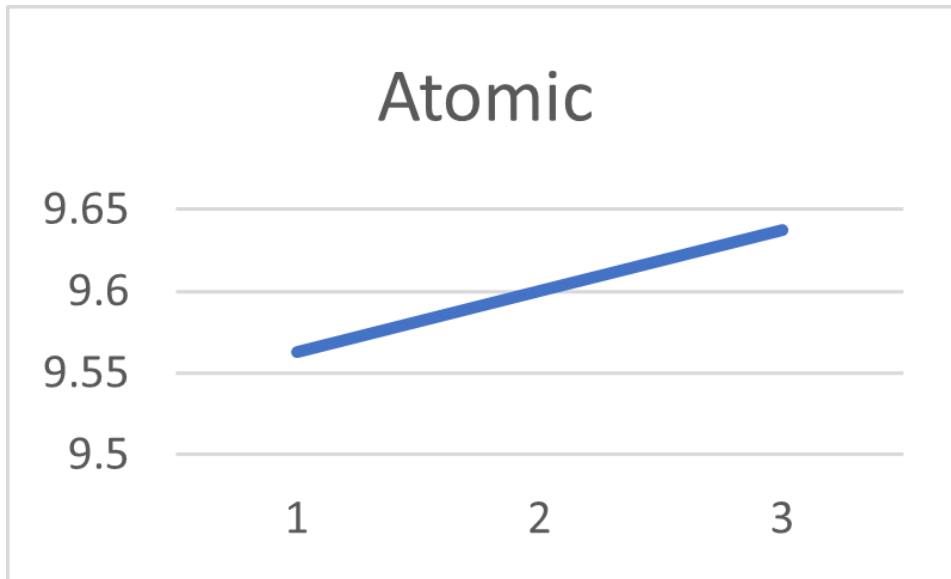
\*1 2 3 represents 2, 4, 8 processes

#### Atomic Version:

```

yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=2 ./histo_atomic uiuc-large.pgm | grep Runtime
Runtime = 59.29 seconds
yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=4 ./histo_atomic uiuc-large.pgm | grep Runtime
Runtime = 59.33 seconds
yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=8 ./histo_atomic uiuc-large.pgm | grep Runtime
Runtime = 59.27 seconds

```



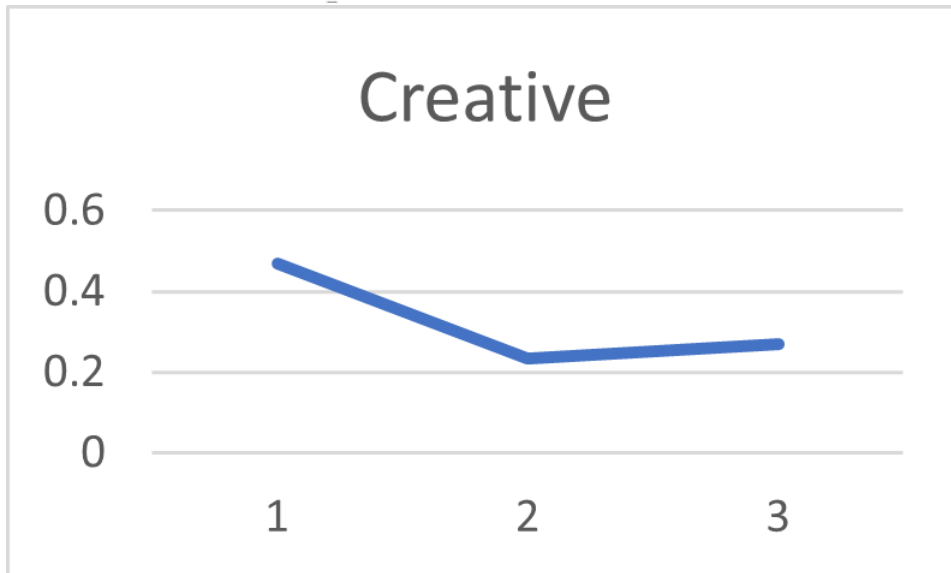
\*1 2 3 represents 2, 4, 8 processes

**Creative Version:**

```

yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=2 ./histo_creative uiuc-large.pgm | grep Runtime
Runtime = 2.91 seconds
yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=4 ./histo_creative uiuc-large.pgm | grep Runtime
Runtime = 1.44 seconds
yl645@vcm-17141:~/openmp/histo$ OMP_NUM_THREADS=8 ./histo_creative uiuc-large.pgm | grep Runtime
Runtime = 1.66 seconds

```



\*1 2 3 represents 2, 4, 8 processes

**Creative Version Demonstration:**

The histo[x] is being updated all the time for the same x, so there will be race condition when executing parallelly. We can use the reduction for the array histo[] to update the value of histo[x] concurrently and sum them up in the end to get correct values for the whole array. To be more

specific, this method will create a private version per process and update the private version concurrently.

**Analysis:**

Speed comparison: creative < original < atomic < lock

The lock version is significantly slower than other versions, this might be due to the cost of communication and the overhead of using locks. Even though I used an array of locks, which means that the granularity of locks is smaller, it still took much longer time than others.

There is not much speed-up for the lock version and atomic version, even worse. But the creative version has a boost when the number of processes increase.

# Problem 2: amgmk

## 1 Code changes

### (1) MATVEC

- Source file: csr\_matvec.c
- Line number: 172
- Code snippet:

```
else
{
    #pragma omp parallel for default(shared) private(i, j, jj, temp)
    for (i = 0; i < num_rows; i++)
    {
        if ( num_vectors==1 )
        {
            temp = y_data[i];
            for (jj = A_i[i]; jj < A_i[i+1]; jj++)
                temp += A_data[jj] * x_data[A_j[jj]];
            y_data[i] = temp;
        }
        else
            for ( j=0; j<num_vectors; ++j )
            {
                temp = y_data[ j*vecstride_y + i*idxstride_y ];
                for (jj = A_i[i]; jj < A_i[i+1]; jj++)
                {
                    temp += A_data[jj] * x_data[ j*vecstride_x + A_j[jj]*idxstride_x ];
                }
                y_data[ j*vecstride_y + i*idxstride_y ] = temp;
            }
    }
}
```

```

temp = beta / alpha;

if (temp != 1.0)
{
    if (temp == 0.0)
    {
        #pragma omp parallel for private(i)
        for (i = 0; i < num_rows*num_vectors; i++)
            y_data[i] = 0.0;
    }
    else
    {
        for (i = 0; i < num_rows*num_vectors; i++)
            y_data[i] *= temp;
    }
}

```

d. OpenMP directive description

As for the first optimization, I set up an omp parallel and for directive with default as shared variable. And private variables are i, j, jj and temp because these variables have read/write conflict. I didn't change the code other than that.

Similarly, I tried the same thing at the second spot, but slightly different in setting up private variables.

e. How do I know to add parallel code.

I used *perf* report:

Percent	Instruction	Assembly
4.22	1b8:	movslq (%rsi,%rax,4),%rcx
3.18		movsd (%r8,%rcx,8),%xmm0
8.55		mulsd (%rdi,%rax,8),%xmm0
14.62		add \$0x1,%rax
4.10		cmp %rax,%rdx
3.28		addsd %xmm0,%xmm1
36.01		jne 1b8
0.09	1d4:	mov (%rsp),%rax
0.64		movsd %xmm1,(%rax)
2.77		jmpq 153

And from this report, I got to know that the most time-consuming part of code is:

```
temp = y_data[i];
for (jj = A_i[i]; jj < A_i[i+1]; jj++)
    temp += A_data[jj] * x_data[A_j[jj]];
y_data[i] = temp;
```

Thus, I built the OpenMP directive on top of this code.

## (2) Relax

- Source file: relax.c
- Line number: 76
- Code snippet:

```
#pragma omp parallel for default(shared) private(i, jj, res)
for (i = 0; i < n; i++) /* interior points first */
{
    /*-----
    * If diagonal is nonzero, relax point i; otherwise, skip it.
    *-----*/
    if ( A_diag_data[A_diag_i[i]] != 0.0)
    {
        res = f_data[i];
        for (jj = A_diag_i[i]+1; jj < A_diag_i[i+1]; jj++)
        {
            // ii = A_diag_j[jj];
            res -= A_diag_data[jj] * u_data[A_diag_j[jj]];
        }
        u_data[i] = res / A_diag_data[A_diag_i[i]];
    }
}

return(relax_error);
```

- OpenMP directive description

I set up an omp parallel and for directive with default as shared variable. And private variables are i, jj and res because these variables have read/write conflict. I also removed variable ii in order to decrease a private storage overhead.

## (3) Axy

- Source file: vector.c
- Line number: 383
- Code snippet:

```

#pragma omp parallel for default(shared) private(i)
for (i = 0; i < size; i++)
    y_data[i] += alpha * x_data[i];
return ierr;
}

```

d. OpenMP directive description

I set up an omp parallel and for directive with default as shared variable. And private variable is only i. I didn't change the code other than that.

## 2 Performance Summary

(1) Sequential.

Run `./AMGMk` 5 times, get averagely:

Total Wall time = 2.86 seconds.

(2) Parallelize with 1 thread

Run `OMP_NUM_THREADS=1 ./AMGMk` 5 times, get averagely,

Total Wall time = 2.87 seconds

(3) Parallelize with 2 thread

Run `OMP_NUM_THREADS=2 ./AMGMk` 5 times, get averagely,

Total Wall time = 1.44 seconds

(4) Parallelize with 4 thread

Run `OMP_NUM_THREADS=4 ./AMGMk` 5 times, get averagely,

Total Wall time = 0.74 seconds

(5) Parallelize with 8 thread

Run `OMP_NUM_THREADS=8 ./AMGMk` 5 times, get averagely,

Total Wall time = 0.42 seconds

In conclusion, as the number of threads goes up, the performance increases, but not strictly proportionally.